FEDERAL AI MONITORING PLAYBOOK

SYNERGIST
TECHNOLOGY

Bahar Ghahremani Bahar Ghahremani
bahar@newworldangels.com bahar@newworldangels.com
Mon Mar 16 10:13:34 2026 Mon Mar 16 10:13:34 2026

Table of Contents

Introduction **2**

Executive Summary of Current Requirements **2**

Governance Structure and Leadership Requirements **3**

Chief AI Officer Mandate 3

AI Governance Boards and Interagency Coordination 4

High-Impact AI: Definition and Risk Management Requirements **4**

Defining High-Impact AI Systems 4

Seven Minimum Risk Management Practices 5

Testing, Evaluation, and Red-Teaming Requirements **6**

Federal AI Testing Standards 6

Red-Teaming and Adversarial Testing 6

Specific Testing Requirements for AI Systems 7

Procurement Requirements and Vendor Obligations **8**

OMB M-25-22: AI Procurement Guidance 8

Vendor Documentation and Transparency Requirements 9

Contract Vehicles and Acquisition Pathways **9**

National Security AI Requirements **10**

National Security Memorandum 25 10

AISI National Security Testing 10

Department of Defense Responsible AI Strategy 11

Solicitations and Requests for Information **12**

Recent Federal AI Solicitations 12

Defense and Intelligence AI Acquisition **12**

Transparency, Disclosure, and Public Reporting **13**

Annual AI Use Case Inventory 13

Public Disclosure of High-Impact Determinations and Waivers 14

Emerging State and Platform Disclosure Requirements 14

Implementation Challenges and Compliance Strategies **15**

Barriers to AI Adoption and Compliance 15

Agency Compliance Strategies 16

Future Outlook and Evolving Requirements **16**

Emerging Policy Developments 16

Technology-Specific Guidance Development 17

International Harmonization Efforts 18

Conclusion **18**

Introduction

The federal government has established an extensive and evolving framework for artificial intelligence monitoring, governance, and procurement that reflects both innovation priorities and risk management imperatives. As of December 2025, agencies face a complex landscape of requirements driven by executive orders, OMB memoranda, national security directives, and legislative mandates that collectively shape how AI systems are developed, procured, tested, and deployed across government operations. Drawing on this landscape, Synergist Technology synthesized primary federal source documents, implementation guidance, and emerging best practices into this playbook to give agencies and their partners a practical, structured view of what federal AI monitoring now requires in real-world practice.

Executive Summary of Current Requirements

The current federal AI monitoring framework is primarily governed by OMB Memorandum M-25-21 (issued April 3, 2025) titled "Accelerating Federal Use of AI through Innovation, Governance, and Public Trust," and its companion OMB Memorandum M-25-22 focused on AI procurement. These directives build upon Executive Order 14110 from October 2023, which established foundational requirements for safe and trustworthy AI development, and Executive Order 14179 from January 2025, which shifted emphasis toward removing regulatory barriers to American AI leadership. Together, these policies create a three-pillar framework centered on innovation, governance, and public accountability.

Federal agencies are now required to appoint Chief AI Officers within 60 days, establish AI Governance Boards within 90 days, and develop comprehensive AI strategies within 180 days. For "high-impact AI", systems that affect safety, civil rights, or access to critical services, agencies must implement seven minimum risk management practices within 365 days, including pre-deployment testing, impact assessments, ongoing monitoring, human oversight, and public consultation. The framework mandates annual public inventories of AI use cases, transparent disclosure of high-impact determinations and waivers, and continuous performance monitoring throughout AI system lifecycles.

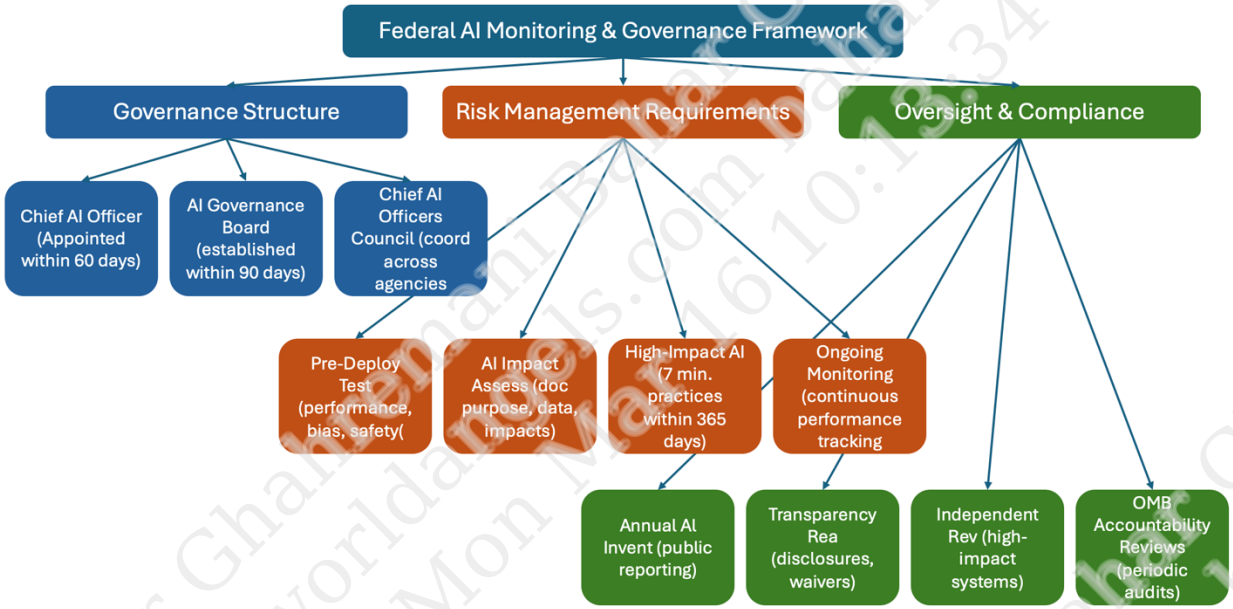

Federal AI Monitoring and Governance Framework - A hierarchical overview of governance structures, risk management requirements, and oversight mechanisms mandated by OMB M-25-21 and related federal AI policies.

Governance Structure and Leadership Requirements

Chief AI Officer Mandate

Every federal agency must now designate a Chief AI Officer (CAIO) who serves as the primary authority for coordinating AI use, managing risks, and ensuring compliance with federal policies. The CAIO role represents a fundamental shift toward centralized accountability, with responsibilities that include advising senior leadership on AI initiatives, maintaining oversight of all agency AI activities, developing compliance strategies, and managing the AI use case inventory.

The position has evolved significantly from its initial conception. In July 2025, Gregory Barbaccia, the Federal Chief Information Officer, was designated as the first Federal Chief AI Officer, effectively creating a government-wide coordination role that sits atop the individual agency CAIOs. This Federal CAIO position, while not explicitly mandated in legislation or executive orders, emerged organically to lead the Chief AI Officers Council and coordinate implementation of the AI Action Plan across all federal agencies.

At the agency level, CAIOs possess substantial authority, including the power to waive specific risk management requirements when compliance would "increase risks to safety or rights overall or would create an unacceptable impediment to critical agency operations". However, these waivers must be tracked, publicly disclosed, recertified

annually, and reported to OMB within 30 days of any changes. The CAIO must also certify pilot programs for high-impact AI and ensure that affected individuals can opt in or out of participation with sufficient notice.

AI Governance Boards and Interagency Coordination

Within 90 days of the April 2025 OMB memo, agencies must establish AI Governance Boards composed of cross-functional leadership including legal, privacy, security, procurement, and program officials. These boards serve as the primary oversight mechanism for AI strategy implementation, ensuring alignment with agency missions while managing technical, legal, and ethical risks.

The governance structure extends beyond individual agencies through the Chief AI Officers Council, which OMB convenes and chairs to coordinate AI development and use across federal programs and operations. This council is tasked with developing shared templates and technical resources, promoting best practices, and facilitating interagency collaboration on AI adoption challenges. The council automatically sunsets five years after creation unless OMB extends its charter.

High-Impact AI: Definition and Risk Management Requirements

Defining High-Impact AI Systems

The cornerstone of federal AI monitoring requirements centers on the concept of **"high-impact AI"**, artificial intelligence systems that make or substantially contribute to consequential decisions affecting individuals' safety, rights, or access to critical resources. According to OMB M-25-21, an AI system qualifies as high-impact if it influences decisions related to:

- Civil rights, civil liberties, or privacy protections
- Access to education, housing, insurance, credit, employment, and government programs
- Access to critical infrastructure or public safety systems
- Strategic assets or national security resources

This definition encompasses both "safety-impacting AI" (systems posing risks to human safety or community well-being) and "rights-impacting AI" (systems with potential to significantly affect civil rights, civil liberties, or privacy). The framework presumes certain AI applications are high-impact unless an agency CAIO, in coordination with relevant officials, formally determines otherwise through a documented process.

Seven Minimum Risk Management Practices

For high-impact AI systems, agencies must implement seven comprehensive risk management practices:

1. Pre-Deployment Testing: Agencies must conduct rigorous testing before operational use to assess AI performance against intended purposes, evaluate potential for biased or discriminatory outputs, and verify safety under realistic conditions. This testing requirement extends beyond technical validation to include assessment of how the system performs across diverse populations and scenarios.

2. AI Impact Assessment: Before deployment, agencies must complete detailed impact assessments documenting the AI's intended purpose and expected benefits, quality and appropriateness of training data and model capabilities, potential positive and negative impacts on affected populations, reassessment scheduling and procedures, cost-benefit analysis, and results of independent reviews. These assessments function similarly to privacy impact assessments but with broader scope encompassing safety, equity, and effectiveness considerations.

3. Ongoing Monitoring: Continuous monitoring requirements mandate that agencies track AI performance metrics, detect degradation or drift in model accuracy, identify adverse impacts on affected populations, and maintain audit logs of system use and outputs. This monitoring must continue throughout the AI system's operational lifecycle, not just during initial deployment.

4. Human Training and Assessment: Agencies must ensure that personnel operating or overseeing AI systems receive adequate training on system capabilities, limitations, and appropriate use contexts. This includes technical training for operators and awareness training for managers and oversight personnel.

5. Human Oversight and Intervention: High-impact AI requires meaningful human involvement in decision-making processes, with clear mechanisms for human intervention when AI outputs require review or correction. The framework explicitly prohibits full automation of consequential decisions without human judgment.

6. Remedies and Appeals: Agencies must provide consistent processes for individuals to challenge or appeal decisions influenced by AI systems. This includes clear explanation of how AI contributed to decisions and accessible mechanisms for human review.

7. Public Consultation: Agencies must consult with and incorporate feedback from end users, affected communities, and relevant stakeholders during AI development and

deployment. This requirement recognizes that those most affected by AI systems often have critical insights into potential harms and necessary safeguards.

Agencies have 365 days from the April 2025 OMB memo to document implementation of these practices and must be prepared to report compliance to OMB through periodic accountability reviews.

Testing, Evaluation, and Red-Teaming Requirements

Federal AI Testing Standards

The federal government has established comprehensive testing requirements that span multiple dimensions of AI system performance, safety, and security. At the forefront of these efforts is the National Institute of Standards and Technology (NIST), which operates through its Center for AI Standards and Innovation (CAISI) to develop measurement science, standards, and evaluation frameworks for trustworthy AI.

NIST's testing framework is built around the AI Risk Management Framework (AI RMF), released in January 2023 and continuously updated to address emerging technologies like generative AI. The framework provides four core functions: **Govern, Map, Measure, and Manage** that guide organizations through identifying, assessing, and mitigating AI risks throughout system lifecycles. While the AI RMF is voluntary for private sector entities, it has become the de facto standard referenced in federal AI policies, including OMB M-25-21's requirement that agencies follow frameworks "consistent with" NIST guidance.

The U.S. AI Safety Institute (AISI) within NIST serves as the primary federal entity for conducting pre-deployment testing of frontier AI models. Under National Security Memorandum 25, AISI is tasked with pursuing voluntary preliminary testing of at least two frontier AI models prior to public release to evaluate capabilities that might pose national security threats. This includes assessing models' capacity to aid offensive cyber operations, accelerate biological or chemical weapon development, autonomously carry out malicious behavior, or automate deployment of other dangerous AI systems.

Red-Teaming and Adversarial Testing

AI red-teaming, the use of adversarial techniques to probe AI systems for vulnerabilities, biases, and failure modes, has become a mandated practice for federal AI procurement and deployment. The practice involves simulating attacks by malicious actors to identify exploitable weaknesses before systems reach operational environments.

OMB M-25-22 requires that agencies procuring enterprise-wide generative AI include contractual provisions for vendors to provide detailed documentation of red-teaming

results. This documentation must be "sufficiently detailed so that agencies can understand the underlying technical and analytical basis for the conclusions of the evaluations, testing, or red-teaming, and reproduce the results where appropriate". The requirement applies both to testing conducted by vendors and to independent third-party evaluations, including those performed by NIST or specialized contractors.

Recent guidance from the Mitre Corporation emphasizes that red-teaming must occur continuously throughout AI system lifecycles—during development, at deployment, and during ongoing operations. Mitre recommends that red-teaming be performed by independent parties on high-risk AI systems prior to executive branch acquisition, with appropriate government agencies like NIST developing standards that mandate red-teaming for government contractors.

Specific Testing Requirements for AI Systems

Federal agencies and contractors face multiple specialized testing requirements depending on AI system type and use case:

Bias and Discrimination Testing: All AI systems used in employment, lending, housing, or other contexts covered by civil rights laws must undergo testing for discriminatory impacts. The Office of Federal Contract Compliance Programs (OFCCP) actively investigates federal contractors' use of AI in employment decisions, requiring validation strategies that meet the Uniform Guidelines on Employee Selection Procedures. Contractors must verify that AI systems perform consistently across demographic groups and do not produce disparate impacts on protected classes.

Biometric System Requirements: AI-based biometric systems face heightened scrutiny and must meet specific technical standards. Agencies must contractually require that biometric AI systems use configurable minimum similarity thresholds for candidate results, employ minimum quality criteria for input data, return candidate matches with similarity scores in one-to-many searches, and maintain comprehensive audit logs capturing input and output data with appropriate privacy protections.

Cybersecurity and Resilience Testing: The National Security Agency, through its AI Security Center, conducts classified testing of AI models' capacity to detect, generate, and exacerbate offensive cyber threats. This testing assesses whether AI systems, if misused, could accelerate offensive cyber operations or automate exploitation of vulnerabilities. For AI systems deployed on national security systems, the NSA issues minimum cybersecurity guidance that must be incorporated into agency AI governance.

Generative AI Evaluation: Agencies deploying generative AI face specific testing requirements around hallucination detection, harmful content generation, and

environmental impact mitigation. The General Services Administration's USAi.gov platform enables agencies to test leading AI models, including those from OpenAI, Amazon, Anthropic, Google, Meta, and Microsoft, before procurement, with built-in bias scoring and performance telemetry. Each model undergoes three evaluations before inclusion: safety assessment (analyzing hate speech and harmful outputs), performance evaluation (question-answering accuracy), and red-teaming for durability.

NIST is actively developing comprehensive standards for AI testing, evaluation, verification, and validation (TEVV) through its zero drafts project. This initiative seeks stakeholder input on creating flexible frameworks that define consistent terminology, lifecycle stages, and guiding principles for AI assessment. The framework emphasizes practical feasibility, reliability, documentation requirements, and risk-based approaches tailored to system use contexts.

Procurement Requirements and Vendor Obligations

OMB M-25-22: AI Procurement Guidance

OMB Memorandum M-25-22, released alongside M-25-21 in April 2025, establishes comprehensive requirements for how federal agencies acquire AI systems and services. The memo focuses on three overarching objectives: fostering a competitive American AI marketplace, safeguarding taxpayer resources through performance tracking and risk management, and promoting effective acquisition through cross-functional engagement.

Priority for American-Made AI: The procurement guidance explicitly prioritizes acquisition of AI systems developed in the United States. Agencies must "maximize use of American AI" and adopt procurement practices that "preference interoperable AI products and services" to sustain a robust federal AI marketplace. This priority creates competitive advantages for contractors with U.S.-developed AI tools and potentially disadvantages foreign AI providers, particularly those from adversary nations.

Intellectual Property and Data Rights: One of the most significant procurement requirements addresses the treatment of government data and intellectual property rights in AI contracts. Agencies must establish policies and processes that:

- **Scope licensing and IP rights** based on intended AI use to avoid vendor lock-in
- **Ensure long-term access** to components necessary to operate and monitor AI systems
- **Provide clear guidance** on handling, access, and use of agency data, limiting collection and retention to what is reasonably necessary

- **Permanently prohibit** vendors from using non-public agency input and output data to train publicly or commercially available AI algorithms without explicit consent

These provisions respond to longstanding concerns about government agencies becoming dependent on proprietary AI systems that cannot be independently audited, modified, or migrated to alternative vendors.

Vendor Documentation and Transparency Requirements

Contractors providing AI systems or services to federal agencies face extensive documentation obligations designed to facilitate transparency, explainability, and performance tracking. OMB M-25-22 requires agencies to obtain from vendors:

System Capabilities and Limitations: Documentation of known AI capabilities, constraints, and guidelines for intended use, including scenarios where the system may perform poorly or unreliably.

Data Provenance: Detailed documentation of training data sources, including information about data collection methods, preprocessing steps, and potential biases in underlying datasets. This requirement extends to documenting "components necessary to operate and monitor the AI system" for as long as necessary.

Testing and Validation Results: Evidence of ongoing testing and validation of AI model performance, effectiveness of AI offerings, and associated risk management measures. For generative AI, this includes red-teaming results and steps taken to mitigate issues discovered during evaluation.

Performance Metrics and Monitoring: Systems that enable agencies to track AI performance and effectiveness over time, including capabilities for detecting model drift, performance degradation, or emergent harmful behaviors.

Contract Vehicles and Acquisition Pathways

Federal agencies employ multiple acquisition vehicles tailored to different AI procurement scenarios:

GSA Multiple Award Schedules: In August 2025, the General Services Administration added major generative AI providers including OpenAI, Google, and Anthropic to its Multiple Award Schedule, expanding federal agency access to commercial AI tools under pre-negotiated contract terms. This addition streamlines procurement but requires that schedule holders ensure their offerings meet federal AI requirements, including bias testing, transparency provisions, and data handling restrictions.

Department of Defense Tradewinds Marketplace: The DoD's Tradewinds platform serves as a specialized acquisition ecosystem enabling rapid sourcing of AI, data, and digital solutions for defense applications. Tradewinds supports various contract types including Other Transaction Authority (OTA) agreements, which enable non-traditional defense contractors and innovative technology providers to engage more easily with DoD. The platform includes a marketplace of vetted AI solutions, a collaboration hub for co-creating mission-aligned capabilities, and streamlined contracting pathways that reduce administrative burden.

Performance-Based Contracting: OMB M-25-22 strongly encourages agencies to use performance-based acquisition techniques that focus on outcomes rather than prescriptive technical specifications. This approach allows agencies to "understand and assess vendor claims about proposed AI systems prior to contract award, acquire AI capabilities that address their needs, and perform post-award monitoring".

National Security AI Requirements

National Security Memorandum 25

On October 24, 2024, President Biden issued National Security Memorandum 25 (NSM-25), the first comprehensive articulation of U.S. national security strategy and policy toward artificial intelligence. The memorandum fulfills a directive from Executive Order 14110 and provides detailed guidance on appropriately harnessing AI models and technologies in national security systems while protecting human rights, civil liberties, and safety.

NSM-25 establishes four primary objectives:

1. **Strengthen and protect the U.S. AI ecosystem** by identifying foreign intelligence threats to AI supply chains and critical enabling sectors
2. **Improve safety, security, and trustworthiness** of AI systems developed and used in the United States through enhanced testing and evaluation
3. **Enhance responsible and effective adoption** of AI across national security agencies
4. **Minimize AI misuse worldwide** through international engagement and export controls

AISI National Security Testing

Under NSM-25, the U.S. AI Safety Institute serves as the "primary point of contact" for private sector AI companies on testing and evaluation activities related to national security concerns. AISI must pursue voluntary preliminary testing of at least two frontier AI models within 180 days to evaluate national security threats, including capabilities to aid offensive

cyber operations, accelerate weapons development, or autonomously execute malicious behaviors.

The AI Security Center at the National Security Agency collaborates with AISI to develop capabilities for rapid systematic classified testing of AI models' capacity to detect, generate, and exacerbate offensive cyber threats. This testing infrastructure assesses the degree to which AI systems, if misused, could accelerate cyberattacks or enable adversaries to exploit vulnerabilities in critical systems.

Department of Defense Responsible AI Strategy

The Department of Defense has implemented its own comprehensive Responsible AI Strategy and Implementation Pathway, first adopted in June 2022 and continuously updated. The DoD strategy is organized around six foundational tenets:

1. RAI Governance: Modernized governance structures enabling continuous oversight of AI use in context-appropriate ways.
2. Warfighter Trust: Ensuring operators achieve justified confidence in AI through technological familiarity and proficiency.
3. AI Product and Acquisition Lifecycle: Exercising appropriate care throughout development, procurement, testing, and deployment to address risks from the outset.
4. Requirements Validation: Using formal processes to ensure AI capabilities align with operational needs while addressing relevant risks.
5. Responsible AI Ecosystem: Building domestic and international partnerships to promote shared understanding of responsible AI design and use.
6. AI Workforce: Ensuring all DoD AI personnel possess appropriate understanding of technology, development processes, and operational methods.

The DoD's five AI Ethical Principles: Responsible, Equitable, Traceable, Reliable, and Governable, apply to all defense AI capabilities, encompassing both combat and non-combat applications. Government contractors developing AI for defense applications must align with these principles, which require comprehensive documentation of data sources and model development, robust bias detection and mitigation, regular security assessments against NIST standards, and governance structures that maintain alignment with DoD principles.

Solicitations and Requests for Information

Recent Federal AI Solicitations

Federal agencies have issued numerous requests for information (RFIs) and solicitations to understand market capabilities and acquire AI-related services:

GSA AI-Powered Procurement System (August 2025): The General Services Administration issued a broad RFI seeking feedback on creating a single, end-to-end integrated procurement ecosystem that incorporates artificial intelligence. The initiative aims to consolidate systems, provide data-driven insights, and modernize how government agencies and suppliers collaborate. GSA specifically seeks input from organizations experienced in IT, data architecture, AI/machine learning, data analytics, and user experience design, with particular emphasis on hearing from small businesses and startups.

Department of Energy AI Infrastructure RFI (April 2025): DOE released an RFI to assess industry interest in developing, operating, and maintaining AI infrastructure on select DOE-owned or managed lands. The initiative seeks to enable construction of AI data centers at DOE sites to begin by end of 2025, with target commencement of operations by end of 2027. The RFI addresses potential development approaches, financing mechanisms, and integration with existing DOE facilities that already have extensive site characterization and permitting for energy infrastructure.

Treasury Bureau of Fiscal Service Records Management RFI (April 2025): The Treasury Department's Bureau of Fiscal Service issued an RFI seeking commercial-off-the-shelf (COTS) records management applications using artificial intelligence technology. The RFI aims to identify potential sources capable of and interested in providing AI-powered records management solutions for federal use.

OPM AI-Enabled Assessment Platform (October 2025): Through GSA Assisted Acquisition Services, the Office of Personnel Management seeks a scalable commercial online assessment platform capable of supporting millions of federal job applicants annually. The $750 million initiative explores leveraging artificial intelligence and machine learning in developing and scoring assessments, preventing applicant misuse of AI, and connecting different assessment systems across government.

Defense and Intelligence AI Acquisition

DoD Tradewinds Marketplace: The Chief Digital and Artificial Intelligence Office operate Tradewinds as an ongoing acquisition vehicle for rapid AI procurement. The platform eliminates traditional procurement friction by enabling streamlined solicitations, down-

selects, and awards, particularly through Other Transaction Authority agreements that facilitate engagement with non-traditional contractors.

AI Contracts Market Growth: Federal agencies are projected to spend over $3.3 billion on artificial intelligence in fiscal year 2025, with total AI spending from FY 2022 to 2024 reaching $5.6 billion—a 6% growth trajectory. The Department of Defense accounts for 72% of all AI contract obligations, totaling $4.0 billion during this period, while small businesses captured 35% of total obligations with 34% growth.

Recent Major Awards: In September 2025, the Pentagon awarded contracts totaling up to $800 million to four technology companies for developing military applications of general-purpose AI models. The Army announced plans to purchase up to $10 billion in software capabilities from Palantir for AI-enabled battlefield decision-making.

Transparency, Disclosure, and Public Reporting

Annual AI Use Case Inventory

Executive Order 14110 and the AI in Government Act of 2020 require federal agencies to conduct annual inventories of their AI use cases and make these inventories publicly accessible. The inventory requirement aims to increase transparency about how government uses AI, enable oversight and accountability, and facilitate knowledge sharing across agencies.

Inventory Content Requirements: Agencies must inventory planned, new, and existing AI use cases, retaining retired or decommissioned cases from prior inventories rather than removing them. For each use case, agencies report:

- **System description** and intended purpose
- **Performance domains** and capabilities employed
- **Data sources** and training data characteristics
- **Rights and safety impacts** (whether the system affects civil rights, civil liberties, privacy, or safety)
- **Development status** and operational deployment stage
- **Compliance status** with risk management practices

Submission and Publication Deadlines: Agencies must submit inventories to OMB by December 16 annually and subsequently publish machine-readable CSV files on their agency websites at [agency.gov]/ai. For 2025, compilation deadlines were set for November 4 with publication by December 2. The Department of Justice's 2024 inventory included

241 entries—a 1,507% increase from 2023—demonstrating rapid AI adoption across government.

Exemptions and Aggregated Reporting: Certain AI use cases are exempt from individual inventory reporting, including those within the Department of Defense, Intelligence Community, and systems whose disclosure would be inconsistent with exemptions in 5 U.S.C. § 552 (Freedom of Information Act). However, agencies must still collect information on exempted cases internally and report aggregate metrics to OMB and publicly, including total number of use cases, percentage classified as rights-impacting or safety-impacting, and number of determinations, extensions, and waivers granted.

Public Disclosure of High-Impact Determinations and Waivers

OMB M-25-21 requires unprecedented transparency regarding agencies' treatment of high-impact AI systems. By December 1 annually, agency CAIOs must recertify and publicly report the full list of active determinations that specific use cases do not qualify as high-impact despite meeting presumptive criteria. These determinations must include justification for why systems that appear to affect safety or rights do not meet the formal definitions.

Similarly, agencies must publicly disclose waivers granted for high-impact AI systems that cannot meet all seven minimum risk management practices. Waivers must be reported to OMB within 30 days of issuance or revocation and recertified annually. The public summaries must describe each determination and waiver along with its justification, ensuring that the public can scrutinize agencies' decisions about when and how to deploy consequential AI systems.

Emerging State and Platform Disclosure Requirements

While federal disclosure requirements focus on government AI use, parallel developments at state and platform levels create additional compliance considerations for AI developers and deployers:

State AI Disclosure Laws: As of late 2025, several states have enacted or proposed legislation requiring disclosure when businesses use AI to interact with consumers. Utah's Artificial Intelligence Policy Act (effective May 2024) requires businesses in regulated occupations to prominently disclose generative AI use at the beginning of communications. California's AI Transparency Act (SB 942, effective January 2026) mandates that covered providers offer free, publicly accessible AI content detection tools and disclosure options. Colorado's AI Act requires disclosure when AI makes consequential decisions affecting employment, education, financial services, healthcare, housing, insurance, or legal services.

Federal AI Disclosure Act: The proposed federal AI Disclosure Act of 2023 (H.R. 3831) would require all generative AI output to include a disclaimer stating, "this output has been generated by artificial intelligence". While the bill has not advanced beyond committee stage, it signals congressional interest in comprehensive disclosure requirements.

Legal Practice Disclosure Requirements: Multiple federal and state courts have implemented local rules requiring attorneys to disclose AI use in legal filings. The Northern District of Texas requires AI disclosure statements on the first page of filings, while Missouri's 20th Judicial Circuit demands disclosure of specific AI tools used. These requirements aim to prevent AI-generated hallucinations from reaching courts while maintaining attorney accountability

Implementation Challenges and Compliance Strategies

Barriers to AI Adoption and Compliance

Federal agencies' 2025 AI compliance plans reveal persistent barriers that complicate both AI adoption and compliance with monitoring requirements. The most commonly cited obstacles include:

Data Access and Quality Issues: Agencies struggle with accessing sufficient high-quality data for training and operating AI systems while maintaining appropriate privacy and security controls. The siloed nature of government data, complex approval processes for data use, and stringent security clearance requirements for classified information significantly slow AI development timelines.

IT Infrastructure Limitations: Existing government IT infrastructure often lacks the computational resources, development environments, and continuous integration capabilities needed for modern AI applications. Agencies cite the need for AI-enabling infrastructure across the full lifecycle, including development, testing, deployment, and continuous monitoring.

AI Talent Shortages: Recruiting and retaining personnel with specialized AI skills remains a critical challenge. The compliance plans emphasize the need for AI workforce development, training programs, and strategies to compete with private sector compensation for technical talent.

Resource Constraints: Budget limitations, procurement complexities, and competition for funding constrain agencies' ability to invest in AI systems, testing infrastructure, and compliance mechanisms.

Agency Compliance Strategies

Agencies are adopting various strategies to meet AI monitoring requirements while managing implementation challenges:

Leveraging GSA Shared Services: Multiple agencies indicate plans to leverage GSA's USAi.gov platform for testing and evaluating AI models before procurement. The Consumer Financial Protection Bureau, Commodity Futures Trading Commission, National Science Foundation, and Securities and Exchange Commission all reference using GSA's vetted AI sources and evaluation frameworks in their compliance plans.

Interagency Collaboration: Agencies are establishing working groups, participating in the Chief AI Officers Council, and engaging in knowledge-sharing forums to develop common approaches and avoid duplicative efforts. The Department of Labor plans to leverage "OneGov deals and shared solutions" while bringing in external consultations through interagency workgroups.

Risk-Based Governance Frameworks: Agencies are implementing risk-based approaches that match approval authority to AI risk levels, enabling lower-level review of low-risk use cases while maintaining rigorous oversight of high-impact systems. The Federal Reserve describes developing a framework that matches approval authority to risk levels, expanding which use cases can be reviewed at lower organizational levels.

Centralized CAIO Oversight: Agencies are establishing centralized review processes where the CAIO office evaluates all AI use cases, provides guidance to programs, and conducts regular audits. The Department of Labor's plan describes requiring agencies to report AI use cases quarterly to the CAIO office for review, with mandatory evidence collection and compliance auditing.

Future Outlook and Evolving Requirements

Emerging Policy Developments

The federal AI monitoring landscape continues to evolve rapidly with new policy initiatives and legislative proposals:

America's AI Action Plan (July 2025): The Trump administration's AI Action Plan signals a strong preference for centralized federal regulation over state-level requirements, with explicit efforts to curtail what the administration characterizes as state overreach. The plan directs the Federal Communications Commission to evaluate whether state AI regulations interfere with its authorities under the Communications Act, potentially asserting federal preemption of certain state laws.

Proposed Federal AI Moratorium on State Laws: Contained within the "One Big Beautiful Bill Act" passed by the House of Representatives is a provision for a 10-year "temporary pause" on states' ability to regulate AI models, AI systems, or automated decision systems. The prohibition would preempt state laws specifically targeting AI unless they remove legal impediments to AI deployment, streamline licensing and permitting, or regulate AI in the same manner as non-AI systems. The federal government may withhold broadband grants from states that enforce AI regulations in violation of this provision.

TEST AI Act: Bipartisan legislation introduced in May 2025 would establish a testbed pilot program through collaboration between NIST and the Department of Energy to develop and refine measurement standards for evaluating AI systems. The bill directs use of National Laboratories' capabilities to establish safeguards, enable safe AI exploration, and protect against misuse.

AI Risk Evaluation Act: Senators Josh Hawley and Richard Blumenthal introduced legislation in September 2025 creating a risk evaluation framework helmed by the Department of Energy for advanced AI systems. The bill requires advanced AI models (trained with computing power greater than 10^26 floating-point operations) to participate in secure testing and evaluation before deployment in interstate or foreign commerce, with penalties of $1 million per day for non-compliant operation.

Technology-Specific Guidance Development

OMB has committed to developing additional specialized playbooks and guidance for specific types of AI systems:

Generative AI Procurement Guidance: OMB M-25-22 notes that additional playbooks focused on procurement of generative AI are under development. These will likely address unique challenges of large language models, including hallucination detection, content filtering, training data transparency, and environmental impact assessment.

AI-Based Biometrics Standards: The procurement memo indicates forthcoming guidance specific to AI-based biometric systems, building on existing requirements for configurable thresholds, quality criteria, and audit logging. This guidance will likely address facial recognition, voice authentication, and behavioral biometrics used in security and access control contexts.

Standard AI Contract Clauses: GSA is directed to establish an online repository for agencies to share AI acquisition information and best practices, including language for standard AI contract clauses and negotiated costs. This repository aims to reduce the burden on individual agencies to develop AI-specific contract terms from scratch.

International Harmonization Efforts

The U.S. AI monitoring framework increasingly intersects with international AI governance initiatives:

EU AI Act Alignment: While the Trump administration has signaled resistance to adopting EU-style comprehensive AI regulation, the practical realities of global commerce mean many U.S. companies must comply with both U.S. and EU requirements. The EU AI Act's risk-based categorization system, mandatory impact assessments for high-risk AI, and transparency requirements share conceptual similarities with U.S. federal approaches, creating potential for harmonization even absent formal policy alignment.

NIST International Leadership: NIST's AI Risk Management Framework and measurement science initiatives have gained international recognition, positioning the U.S. as a standard-setter for voluntary AI governance approaches. The AI Safety Institute Consortium brings together over 280 organizations globally to develop science-based guidelines that inform international standards.

National Security AI Engagement: NSM-25 emphasizes international engagement to influence global AI development trajectories consistent with U.S. values and interests. This includes working with allies on AI safety standards, coordinating export controls on AI-enabling technologies, and countering adversary AI capabilities.

Conclusion

The federal government's requirements for AI monitoring have evolved from nascent policy experiments into a comprehensive governance framework that touches every aspect of how agencies develop, procure, test, and deploy artificial intelligence systems. The current framework, anchored by OMB Memoranda M-25-21 and M-25-22, reflects a delicate balance between accelerating AI innovation to maintain U.S. competitiveness and ensuring robust safeguards for safety, civil rights, and public accountability.

For contractors and vendors seeking to engage in AI-related federal procurement, success requires not only technical capabilities but also deep understanding of governance requirements, testing standards, transparency obligations, and risk management practices. The emphasis on American-made AI, prohibition on vendor lock-in, documentation requirements, and continuous monitoring creates both opportunities for well-prepared contractors and compliance challenges for those unprepared for the scrutiny that accompanies AI deployment in government contexts.

As the AI policy landscape continues to evolve with new legislation, updated OMB guidance, NIST standards development, and potential federal preemption of state laws,

organizations must maintain adaptive compliance strategies that anticipate rather than merely react to regulatory changes. The trajectory clearly points toward more, not less, federal oversight of AI systems, but with an emphasis on risk-based approaches that enable innovation in lower-risk contexts while maintaining stringent controls for high-impact applications that affect fundamental rights, safety, and security.

The federal AI monitoring framework ultimately serves multiple constituencies: agencies seeking to leverage AI for mission delivery, contractors offering AI solutions to government, individuals whose rights and safety may be affected by AI decisions, and policymakers tasked with ensuring that American AI leadership aligns with democratic values and constitutional protections. Navigating this complex environment requires ongoing attention to policy developments, active engagement with standards-setting bodies like NIST and the Chief AI Officers Council, and commitment to the principles of responsible AI development that transcend specific regulatory requirements to embody broader ethical and societal obligations.

Call to action:

For help implementing this playbook or assessing your agency's AI monitoring readiness, please contact Mark Bortman at mbortman@synergist.technology or (617) 640-3376 to schedule a working session and get started.

Bahar Ghahremani bahar@newworldangels.com Mon Mar 16 10:13:34 2026