Market Guide for AI Governance Platforms

4 November 2025 - ID G00837249 - 34 min read

By: Lauren Kornutick, Sumit Agarwal, Avivah Litan, Svetlana Sicular, Priya Sundararaman, Souparna Palit

Initiatives: Redefine D&A and AI Governance; Establish Robust AI Governance and Organizational Trust; Technologies and Markets; Strategy, Risks and Opportunities

> AI governance platforms is an emerging market that provides the leader responsible for AI governance with central oversight of AI, application of risk management frameworks and execution of necessary controls. Use this research to better understand the product capabilities in this emerging market.

Overview

Key Findings

- Organizations are extending traditional governance, risk management and compliance structures to address AI oversight. To operationalize these requirements, AI governance platforms (AIGPs) have emerged to provide centralized inventory, risk management, policy enforcement and continuous monitoring across the AI life cycle.
- AIGPs are differentiating themselves from adjacent markets like governance, risk and compliance (GRC), AI security, data science and machine learning (DSML), and data and analytics governance platforms (DAGP), providing central oversight of AI risks and associated runtime enforcement processes.
- Increasing regulatory complexity (e.g., EU AI Act) and a surge in audit coverage for AI-related risks are compelling organizations to invest in AIGPs.
- Many vendors are entering the AIGP market, often positioning themselves as offering comprehensive solutions despite only covering parts of the AI life cycle.
- Both market consolidation and growth are expected, but may risk stifling innovation and leaving buyers dissatisfied if acquired products lose fit-for-purpose capabilities.

Recommendations

- Partner with assurance teams and the board to ensure they have updated their corporate governance mechanisms to clarify responsibility and drive accountability across multiple governance teams.
- Assess the capabilities, described in this market, of existing vendors, including asking for vendor roadmaps. Then, if capabilities are missing or inadequate, select vendors in this market who meet current AI governance needs.
- Prioritize vendors that complement existing solutions in your tech stack to enable scalable AI governance across the entire life cycle, not just single parts of the AI governance process. Select platforms matching the critical gaps in your AI governance process, supporting oversight workflows, and integrating with existing compliance and security systems

Strategic Planning Assumption(s)

- Fragmented AI regulation will see four-fold growth and spread to cover 75% of the world's economies, driving $1 billion in total compliance spend by 2030.
- By 2028, enterprises with revenues above $1 billion will use on average ten different GRC software products, up from eight in 2025.
- By 2029, "death by AI" legal claims will have doubled from the previous decade because decision-automation deployments lacked sufficient AI-risk guardrails.
- By 2028, governance technologies will decrease regulatory compliance costs by 20%, enabling 10% more investment in strategic growth initiatives.

Market Definition

Gartner defines AI governance platforms as tools designed to ensure organizations adhere to organization policy, regulations and industry standards across common responsible AI principles. These platforms allow leaders responsible for AI and other technical or business leaders to streamline governance processes organizationwide and serve as a central repository for trust, risk and security controls. They also automate workflow approvals for new AI use cases, applications and to streamline governance processes organization wide. AI governance platforms support a wide range of AI techniques across built, blended, embedded and bring-your-own-AI applications.

AI governance platforms align the execution of governance processes across the organization. These platforms comprehensively automate, manage and report on AI risks and acceptable use adherence in the enterprise across all forms of AI. They serve as management oversight systems in corporate governance and continuously manage, implement and enforce the necessary trust, risk and cybersecurity controls (e.g., data and model guardrails) to demonstrate the organization has implemented and is governing all of the organization's AI (built and blended, embedded and applications/bring your own AI).

These tools facilitate the AI use case, application or agent risk assessment process; enable workflow automation; and streamline information exchange with other second-line oversight functions. These functions could include AI project value and ROI assessments; privacy; cybersecurity; enterprise risk management; and data and analytics (D&A) governance.

A distinct feature of an AI governance platform is the ability to automate continuous monitoring and runtime enforcement controls from governance rules based on corporate policy, regulations, frameworks and standards. Another distinct feature is the ability to capture data and/or metadata from more than one of the following operational governance categories: acceptable use and other organization policy; security runtime inspection and enforcement functions; observability; and data governance. AI governance platforms have policy engines (e.g., prepackaged rules and/or models) to adhere to common regulations (e.g., EU AI Act), frameworks and standards (National Institute of Standards and Technology AI Risk Management Framework [NIST AI RMF], International Organization for Standardization [ISO] 42001) with the option to customize rules for corporate policy and apply enforcement at runtime.

Another key feature of an AI governance platform is that it serves a wide range of AI techniques across built and blended, embedded and bring your own AI in the enterprise. Therefore the platform must also be interoperable across the organization's technology and data stack.

Mandatory Features

AI Inventory/Catalog: Provides a centralized, discoverable registry of all AI use cases and AI services like SaaS services with embedded AI that can support multiple use cases. These use cases include applications, agents and models within the organization, including version history, metadata (purpose, data sources, algorithms), documentation (e.g., model cards or systems cards), ownership, development stage and deployment status.

Risk Management and Regulations: Provides AI features and risk cataloging applied to AI applications, models or use cases. Frameworks to classify, assess and mitigate AI-specific risks (bias, fairness, robustness, etc.) including content libraries that address unique regulations for AI and privacy (e.g., EU AI Act, GDPR) frameworks (e.g., NIST AI RMF) and standards (e.g., ISO 42001) as well as organization acceptable use and responsible AI policies (also called AI ethics policies).

Automated Policy Compliance and Enforcement at Runtime: Provides centralized, automated management and enforcement of AI-specific policies via multiple guardrails, including control validation for AI-specific risks (e.g., bias, data leakage, trust, privacy, security), access controls, use case alignment and other enterprise policies, remediation recommendations and compliance reporting.

Data Usage Mapping: Captures data mapping used by various AI entities and tracks the usage and misuse over time. This may also include ability to track the provenance of training data and interface with data governance platforms to include data lineage, classification, ownership and data observability information. This feature can be captured via the AI governance platform directly or via interoperability with a D&A governance platform.

Evidence Collection: Provides documentation for trust, risk and security assessments, testing and validation results (such as security, bias detection, model), risk and compliance remediation evidence.

Interoperability: Enables different systems, devices, applications or agents to work together exchanging and utilizing information effectively. Integrations could include the D&A governance platform, model observability tools, AI discovery tools, AI cybersecurity platforms (e.g., red teaming, scanning for all entries, prompt-injection checks), and AI project and ROI management tools.

Workflow and Approvals: Enables the automation of routine governance tasks such as new AI use case model, application or agent approval by a governing body; risk and security assessments internal to the organization or with third parties; approvals; testing procedures; and documentation generation. Facilitates communication and coordinated action among diverse stakeholders. Includes structured signoff, attestation and approval requirements.

Audit: Provides comprehensive audit trails of actions taken in the platform and, where applicable, automatically logs all activities related to the AI life cycle.

Common Features

AI Usage Reporting: Automates the generation of standardized documentation for AI models and use (e.g., model cards, datasheets) for auditors, regulators, etc.

Observability: Enables model builders or owners to monitor, understand and diagnose the performance and behavior of AI models, applications and agents in production. This includes various metrics and signals to ensure models are functioning correctly and delivering accurate predictions. For machine learning models, this is sometimes referred to as "explainability."

Business-Friendly User Experience: Enables targeted users to easily navigate and use the tool to complete their tasks without the need to consult with product subject matter experts (SMEs) or technical staff. This could be interpreted at minimum to mean that the majority of users will not revert back to tools such as spreadsheets after using the tool and/or including visualization capabilities in the user experience (UX).

Ease of Implementation: Allows users to quickly adopt a new instance of the tool to support AI governance activities without the need to heavily customize off-the-shelf templates/prebuilt workflows or make changes to the underlying data model. This could also include how-to guides or easy plug-ins to implement the tool.

Market Description

The leader responsible for AI governance will typically address and treat the risk associated with AI and oversee shared controls owned by other operational governance teams (see A Checklist for Managing Risk Across Six Pillars of AI Governance). To do so and better define these processes, they may rely on AI governance frameworks that align directly to the "govern" function of the U.S. NIST AI Risk Management Framework (RMF), [1] ISO IEC 42001, [2] Gartner AI Maturity Model and others.

These governance frameworks are also a prerequisite for compliance with regulations such as the EU AI Act [3] and similar regulatory obligations. To operationalize these complex governance requirements, leaders responsible for AI governance are turning to AI governance platforms (AIGPs) to improve oversight of AI in their enterprise.

Key Use Cases of AI Governance Platforms

Figure 1 shows the core capabilities of AIGPs, emphasizing their ability to provide end-to-end oversight, automated risk management, policy enforcement and continuous monitoring across the AI life cycle. AIGPs combine these functions to deliver comprehensive, scalable and auditable governance for all types of enterprise AI.

Figure 1: AI Governance Platform Market Overview

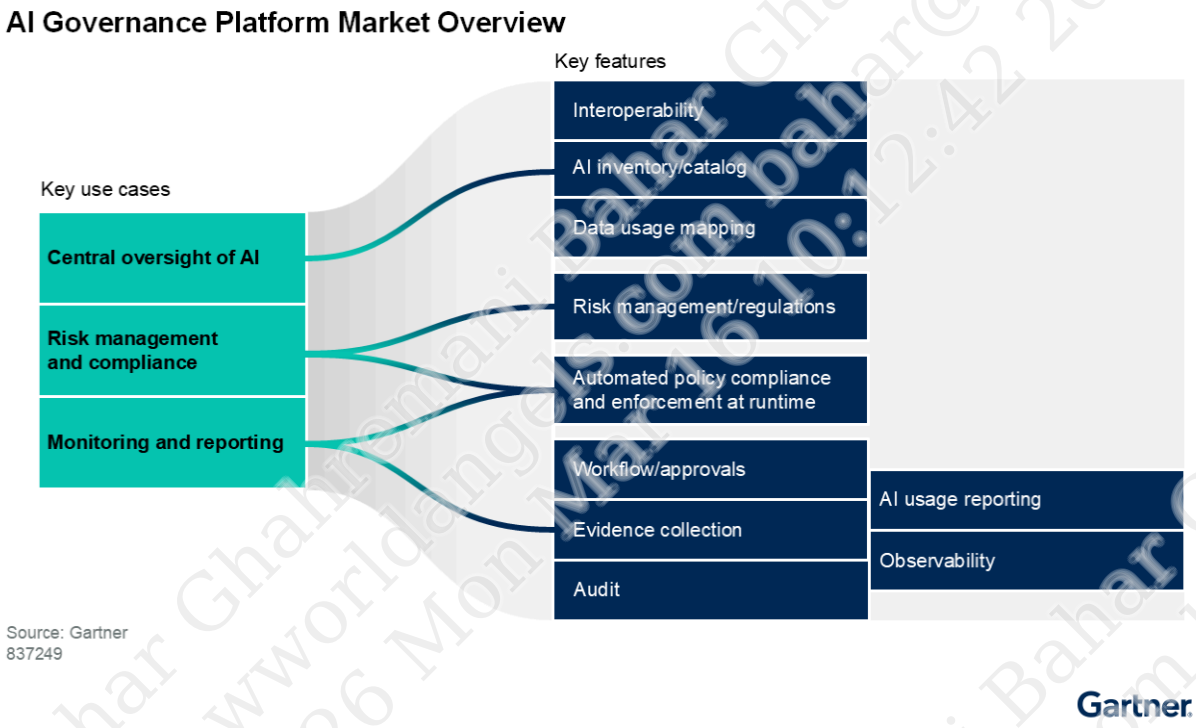

Central Oversight

Central oversight enables governance. It ensures accountability, transparency and appropriate risk management. It also helps track AI systems, monitor compliance, prevent misuse and enable informed decision making, fostering trust and responsible deployment of AI technologies.

Risk Management and Compliance

Given the relationship and direct alignment to risk management frameworks, compliance requirements and organization policy, AI governance platforms enable execution of these processes. Some risk management frameworks and regulations for AI overlap with control requirements that help to secure the organization.

AIGPs typically include some functionality to address the execution and enforcement at runtime of controls, such as acceptable use policies, how data is accessed and moved within and outside the organization into AI models, and model observability.

Monitoring and Reporting

Traditional assurance processes rely on point-in-time audits or periodic testing to validate whether the aforementioned governance frameworks are effective. However, those processes are ineffective as AI is continuously accessing information and making decisions. AI governance platforms have capabilities that detect anomalies, prevent misuse and maintain compliance with regulations and policy.

Monitoring is accomplished through automated audits, real-time data analysis, performance tracking and alert systems. Continuous monitoring and enforcement at runtime is necessary to help identify errors, biases or secondary security threats, allowing timely intervention and reporting to other stakeholders in the organization to mitigate the risk.

Usability

AI governance teams may be led by nontechnical teams that need to implement governance solutions quickly. Products can differentiate by building a business-friendly user experience and prioritizing ease of implementation.

Market Direction

The increasing interest in AIGPs is due to multiple drivers:

- CEOs believe they are entering a new business era, with 77% believing AI will have the most significant impact on their industries over the next three years. [4] As a result, organizations continue to invest heavily in AI technology: 87% of CIO and technology executives reported their enterprise would increase funding for GenAI in 2025, while 84% expected their enterprise to increase funding for traditional AI. [5]
- Internal and external stakeholders mistakenly expect AI systems to be trustworthy, transparent and secure, and as a result of the demands, AI governance is a corporate governance issue. Given the visibility of corporate governance issues, AIGPs are necessary to consolidate information and provide assurance.
- Regulatory complexity is a significant driver for investment in platforms as organizations navigate new complicated regulations governing sovereignty, AI use and data. Regulations can be targeted directly at AI, such as the EU AI Act, U.S. Texas Responsible Artificial Intelligence Governance Act (TRAIGA), and in some cases they overlap into governance for adjacent topics like privacy and security.
- One survey of IT and non-IT professionals found that 78% of respondents were currently developing an agent in their organization, raising the need to evolve governance beyond static policies and incorporate enforcement of controls at runtime. [6]

- There is an increased risk of reputational damage as AI incident reporting is becoming mainstream. The AI Incident Database is dedicated to indexing the collective history of harms or near harms realized in the real world by the deployment of artificial intelligence systems. [7] Multiple or severe incidents can erode trust and cause brand damage.

- There is also a misconception that risk management processes are a barrier to technology adoption. This assumption is incorrect, and organizations with effective AI/data and analytics (D&A) risk management are 12% more advanced in technology adoption stages than organizations with ineffective risk management. [8]

AI Governance Platforms Will Emerge as Specialized Tools

The global AI governance market size accounted for $227.65 million in 2024 and is predicted to increase from $309.01 million in 2025 to approximately $4,834.44 million by 2034, expanding at a CAGR of 35.74% from 2025 to 2034 .[8] The market growth is attributed to the rising adoption of AI across various industries and the need for transparent, accountable decision-making frameworks. [8]

Leaders responsible for AI governance must be aware that vendors in multiple markets (see Figure 2) frequently position themselves as AI governance platforms. However, claims that AI governance can be fully supported in one solution are frequently exaggerated; most only cover parts of the AI life cycle and lack the full-spectrum governance, risk management and policy enforcement required for seamless execution of the AI governance process.

Bahar Ghahremani bahar@newworldangels.com Mon Mar 16 10:12:42 2026

Figure 2: AI Governance Operating Model

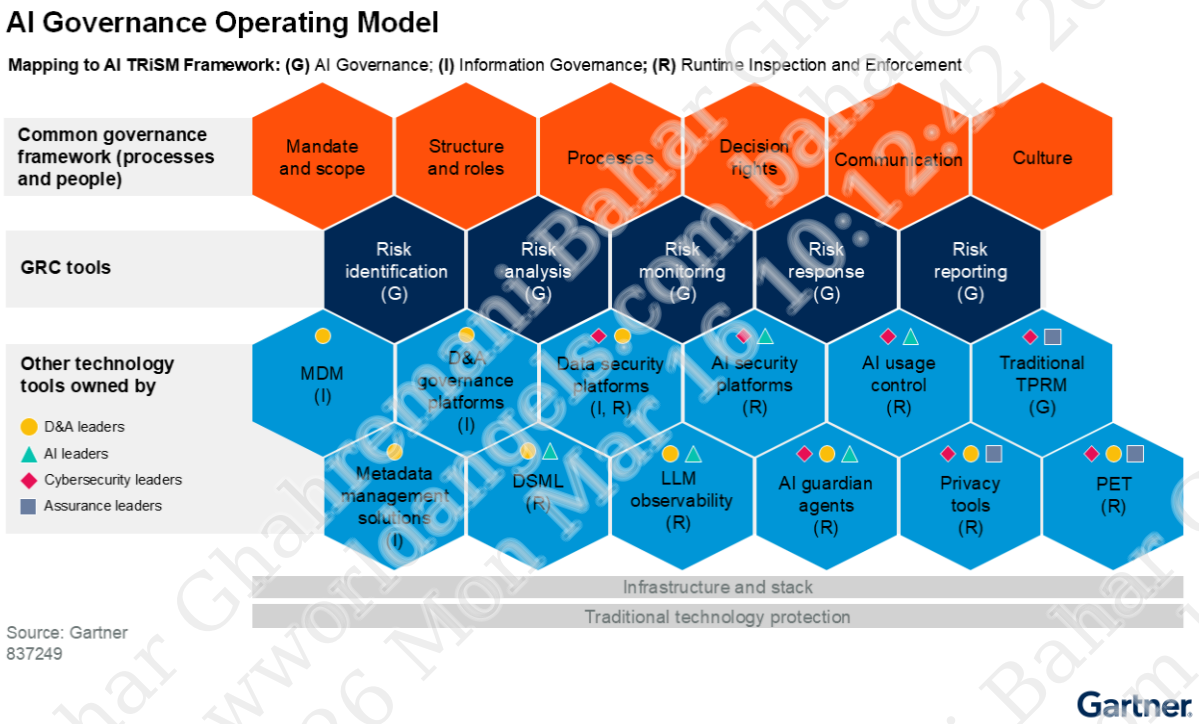

Gartner expects significant market consolidation as buyer requirements become clearer. We have observed that enterprise GRC and AI security platforms will absorb AIGPs and adjacent tools offering runtime inspection and enforcement, targeting buyers in IT GRC and other assurance functions. We expect enterprise vendors will continue to exploit these relationships to push AI governance solutions that may fall short of expectations.

While consolidation is typical for emerging markets and often marks a startup's exit, the reality is far less rosy: postacquisition strategies frequently stifle innovation, resulting in products that quickly lose their edge and fail to meet customer needs. This leaves buyers with uninspired, ill-fitting solutions and little recourse.

Other Alternative and Overlapping Products

As the function for AI governance proliferates in organizations, technology in adjacent markets has evolved to cover specific components of the AI governance process. For that reason, buyers seeking to implement Gartner's AI trust, risk and security management (TRiSM) framework to govern AI need to be aware of what specific capability for AI governance oversight they are looking to fulfill with technology solutions.

Buyers should be mindful of what information they need visibility into to make decisions about AI use and where there may be shared control ownership with existing governance or operational teams. This is especially important for organizations seeking to federate AI governance with existing processes and technologies already in place. Buyers should be aware that this usually translates to a lesser degree of full enterprise oversight by AIGPs, due to a narrower focus of these products on workflows supported primarily by the following teams.

D&A Leaders

- **D&A governance platforms**: D&A governance platforms provide enterprise oversight of D&A governance processes. See Magic Quadrant for Data and Analytics Governance Platforms.
- **Metadata management solutions**: Aggregating solutions that require functions and capabilities to enable enhancement, interoperability and orchestration across systems. See Market Guide for Metadata Management Solutions.
- **Master data management (MDM)**: Provides uniformity, accuracy, stewardship, semantic consistency and accountability of an enterprise's official shared master data assets. See Market Guide for Master Data Management Solutions.

Leaders Responsible for AI

- **DSML platforms** evaluate model predictions and continuously monitor models. See Magic Quadrant for Data Science and Machine Learning Platforms.
- **Large language model (LLM) observability solutions** monitor, analyze and provide actionable insights into the behavior and performance of LLMs. See Innovation Insight: LLM Observability.
- **Guardian agents** are a blend of AI governance and AI runtime inspection and enforcement functions. They oversee, control and secure the operation of AI applications and agents. See Guardians of the Future: How CIOs Can Leverage Guardian Agents for Trustworthy and Secure AI.

Cybersecurity Leaders

- **Data security platforms** combine data security controls with business logic and fine-grained authorization, producing significant gains in efficacy and structured data security. See Market Guide for Data Security Platforms.

- **AI security platforms** are designed primarily to defend against LLM-native security risks. See Use an AI Security Platform to Launch Your AI Security Strategy.
- **AI usage control (AI-UC)** is a technology to discover the use of third-party AI and enforce the security policies of an organization. See Innovation Insight for AI Usage Control.
- **Privacy tools** aim to add compliance controls on personal data as prerequisite to AI-based usage. See State of Privacy: Key Technologies for a Data-Driven Strategy.
- **Privacy-enhancing technologies (PET)** are a collection of forward-looking technologies to protect personal and other sensitive data. See Privacy Program FAQs: Building an Effective Privacy Program.

Assurance Leaders (Legal, Compliance, Enterprise Risk and Audit)

- **GRC tools for assurance leaders** facilitate key areas of the risk management process associated with the diverse enterprise and compliance risks. See Magic Quadrant for Governance, Risk and Compliance Tools, Assurance Leaders.
- **The traditional TPRM technology** market addresses risk management activities of third parties. See Market Guide for Third-Party Risk Management Technology Solutions.

Organizations can use these tools and others to build a strategic technology roadmap for AI governance and determine gaps. However, to address central oversight the data from multiple tools owned by multiple teams would need to be rolled up into broader risk reporting via business intelligence platforms. This will ultimately undermine the effectiveness of a multitool approach, therefore ultimately leading to the selection of an AIGP point solution.

Future Proofing

Buyers must consider how rapidly AI is evolving and the implementation complexity of governance frameworks, and look for the four capabilities below to future-proof their investment:

AI agents: AI agent implementations span a spectrum from automation to autonomous capabilities and are a critical use case for AIGPs. The variability of capabilities and the dynamic nature raise challenges such as reliability, access control, guardrail applications, observability and cost. This creates a complex environment for risk assessment, and regulatory and organizational policy compliance. In addition, AI agents may be deployed on multiple systems and connected using agent-to-agent protocols. AIGPs will need to extend their capabilities to include multisystem AI agents, tools and orchestration to capture the complete spectrum of AI agent capabilities. Enabling transparency and accountability in AI agent-based systems is critical for establishing trust. For this to work successfully, the critical strength and differentiation of AIGP will be the interconnectedness with runtime enforcement tools as a circular or bidirectional feed of information, and effective operational security

Value measurements: AI governance should be an enabler for enterprise AI initiatives. With budget constraints, organizations need to be very selective with use cases, providing business value. AIGPs will need to provide a framework to capture the use-case requirements and the expected business value, or other nonfinancial KPIs (such as usage or hours saved). AIGPs will further need to extend observability capabilities to keep track of the outcomes and the associated value, to enable a single organizational view of the value achievement.

Third-party and embedded AI risk: Oversight and mitigating risk from third parties will also emerge as a critical function. While this could include traditional processes (such as vendor due diligence and risk assessment) in AIGPs, it also extends into a partnership with cybersecurity teams to discover unknown usage of AI by third parties. Therefore, buyers need to look for vendors who have partnerships with AI-UC vendors.

Interoperability: AIGP vendors will need to rethink how they address system interoperability. To truly be effective, they will need to connect to multiple endpoints in the organization's technology stack as well as integrate with gateways. API fees will add up quickly and will render the product useless if charging by the number of connections.

Market Analysis

AI governance tools' value lies in their ability to integrate and synthesize real-time or near real-time information about AI use in the enterprise against predetermined risk tolerances. Therefore, system interoperability with solutions that build AI and third parties that provide enterprise AI are the most critical factors for success.

We expect continued growth through a combination of investment, acquisition and consolidation. We are also observing substantial potential in the AIGP space as:

- AI risk management is prioritized by boards and other assurance teams
- Regulations across the globe are increasing in frequency and complexity
- Data governance and data management requirements are rolled into the AI governance process

Representative Vendors

The vendors listed in this Market Guide do not imply an exhaustive list. This section is intended to provide more understanding of the market and its offerings.

Vendor Selection

The list of vendors in Table 1 is not exhaustive, and it represents vendors that Gartner has identified under the scope of the emerging AIGPs market, specific to the needs of AI governance personas who are primarily responsible for use case approval, risk management, compliance and policy, and where securing the organization is a secondary concern. In addition, representative vendors must support all AI in the enterprise and address governance use cases across multiple AI techniques. (see Note 1). Use this list as a starting point in your exploration of AIGPs.

Bahar Ghahremani bahar@newworldangels.com Mon Mar 16 10:12:42 2026

Table 1: Representative Vendors

(Enlarged table in Appendix)

Vendor	Product Name	Head quarters	Year Founded
Airia	Airia AI Security and Governance	Atlanta, Georgia, United States of America	2023
Cranium	Cranium AI Governance Platform	Short Hills, New Jersey, United States of America	2023
Credo AI	Credo AI Governance Platform	Palo Alto, California, United States of America	2020
Domino Data Lab	Domino Governance	San Francisco, California, United States of America	2014
Dynamo AI	DynamoGuard, DynamoEval	San Francisco, California, United States of America	2021
Fairnow	FairNow AI	McLean, Virginia, United States of America	2023
Holistic AI	Holistic AI	San Jose, California, United States of America	2020
IBM	watsonx.governance	Armonk, New York, United States of America	1911
LatticeFlow AI	AI GO!, AI Insights	Zurich, Switzerland	2020
LogicGate	Risk Cloud	Chicago, Illinois, United States of America	2015
Monitaur	Monitaur Governance Platform	Boston, Massachusetts, United States of America	2019
ModelOp	ModelOp Center	Chicago, Illinois, United States of America	2018
OneTrust	OneTrust AI Governance	Atlanta, Georgia, United States of America	2016
Prodago	Prodago AI Governance	Montreal, Quebec, Canada	2013
Relyance AI	AI Governance Expert	San Francisco, California, United States of America	2020
Saidot	Saidot AI Governance Platform	Helsinki, Finland	2018
Securiti	AI Security & Governance	San Jose, California, United States of America	2019
ServiceNow	AI Control Tower	Santa Clara, California, United States of America	2004
Trustible	Trustible Responsible AI Governance Platform	Arlington, Virginia, United States of America	2023
Truyo	AI Governance and Privacy Automation Platform	Chandler, Arizona, United States of America	2016

Source: Gartner

For additional vendors, please see Tool: Vendor Identification for AI Governance Platforms.

Vendor Profiles

Airia

Airia provides central oversight for AI agents and models, and blends security features with policy-based governance capabilities. The platform collectively addresses oversight of AI via responsible AI guardrails through its Agent Constraints engine, and enterprise security via detailed auditing, runtime monitoring, context-aware permissions, prompt injection protection and secure tool connectivity through the MCP gateway. Airia's governance capabilities include configurable use-case-review workflows, human-in-the-loop oversight, and customizable dashboards for performance and compliance tracking.

Cranium

Cranium combines AI governance and security platforms to provide oversight for securing and verifying AI systems across internal and external environments. The platform provides continuous visibility into AI risk surfaces, regulatory compliance and mitigation of shadow AI and security gaps. Features for AI governance include model discovery, risk assessments and continuous monitoring, and they extend into cybersecurity governance with red-teaming via the Arena module. A differentiating feature is Cranium's AI Card, which can provide insight into the development, deployment and compliance status of all AI systems.

Credo AI

Credo AI offers a centralized AIGP designed to address industry-specific risks, regulatory requirements and organizational policies. The platform enables central oversight and assurance via capabilities that inventory AI use cases, catalog models and datasets, assess and score risks, and generate audit-ready artifacts using Policy Packs aligned with risk management frameworks and regulations such as the NIST AI RMF, ISO 42001 and the EU AI Act. The AIGP provides automated workflows, contextual governance, vendor and third-party risk assessments. The platform's system interoperability and open architecture enable interoperability. Credo AI differentiates by offering an AI Governance Readiness Program via its advisory services.

Domino Data Lab

Domino Data Lab provides an integrated AI governance platform that unifies policy management, model tracking and compliance across all AI assets. Policies are defined, versioned and enforced through automated checks, approvals and gates embedded directly within data science workflows. The AI inventory captures lineage, metadata and compliance status for both ML and GenAI models. Automated documentation, such as model cards, audit trails and evidence notebooks, supports regulatory readiness and transparent collaboration between developers and validators, enabling enterprisewide AI governance.

Dynamo AI

Dynamo AI blends AI governance with security by providing visibility across all AI interactions, supporting both proprietary and third-party models. The platform includes DynamoEval for preproduction testing and evaluation of key risks of AI systems, and DynamoGuard to monitor and prevent compliance violations, data leakage, hallucinations and security jailbreaks in real time. Governance capabilities can be tailored and aligned to organization risk profiles using human-in-the-loop observability and auditability across AI model endpoints, use cases. third-party copilots, and AI-assisted policy writing, thus supporting regulatory compliance. The solution supports interoperability via integration with major cloud platforms and is deployable on-premises.

Fairnow AI

Fairnow AI centralizes AI governance processes by automating risk management and regulatory compliance. The tool enables a centralized AI inventory, streamlined vendor evaluations and collaborative workflows, enabling efficient governance. The platform supports frameworks such as ISO 42001, NIST AI RMF and regulations like the EU AI Act. It provides real-time alerts, risk flagging and compliance guidance. Fairnow automates bias testing, transparency documentation, and aspects of compliance management through intelligence and agentic automation capabilities, integrating with DSML and GRC tools for scalable oversight. See Note 3 for additional information.

Holistic AI

Holistic AI's AIGP solution supports AI governance via risk management, regulatory compliance and AI life cycle oversight. The product includes automated discovery, risk scoring, agents and monitoring that span bias, privacy and robustness. The platform supports global regulations including the EU AI Act and frameworks like ISO 42001, and allows customizable rulebooks and framework builders. Structured governance is enabled through intake forms, stakeholder signoffs and continuous monitoring. Holistic AI extends into cybersecurity governance and includes red-teaming tools, audit-grade documentation and real-time alerts, integrating with enterprise tools like Jira and ServiceNow.

IBM

IBM's watsonx.governance provides an integrated solution for managing AI life cycle risks. The platform centralizes model documentation, automates metadata capture and offers dynamic dashboards for compliance and risk tracking. IBM supports transparency, explainability and bias mitigation, with integrated tools for risk management and retraining recommendations. The solution supports agentic AI governance (AgentOps) and third-party integrations, providing oversight for generative and autonomous systems. IBM's governance framework includes a responsible Technology Board and Office of Privacy and Responsible Technology for alignment with industry standards and global regulations, such as the EU AI Act. Guardium AI Security extends security capabilities.

LatticeFlow AI

LatticeFlow AI integrates governance frameworks with technical model assessments, supporting compliance with regulations such as the EU AI Act. It supports generative AI, computer vision and structured data models. The platform offers templates to create and manage AI risk frameworks, map regulatory and governance requirements to technical controls, and perform evaluations related to bias, security, robustness and performance. It includes features for risk scoring, regulatory mapping and generating technical documentation for audits.

LogicGate

LogicGate's AIGP, part of their Risk Cloud suite, leverages Spark AI to deliver configurable workflows for risk identification, control creation, assessments and anomaly detection. AI features include autofilling risks and controls from documents, recommending record linkages and dynamic gap analysis. The platform supports integration with cyber, compliance and third-party risk programs. Risk Cloud also offers single-tenant deployment options and access-controlled environments for heightened data privacy or regulatory requirements.

Monitaur

Monitaur enables oversight for built and purchased AI across predictive, generative and agentic systems. It supports AI inventory, regulatory and standards management, governance workflows and automated policy enforcement. Govern enables stakeholder collaboration with a business-friendly user experience. Automate continuously monitors and validates data and AI systems to satisfy governance policies and provide assurance. It monitors flags and risks like bias, drift and performance issues and suggests remediations. Integration with enterprise risk, data and modeling platforms is supported with libraries and an API-first approach. The platform differentiates itself with dedicated support, expert reviewer services, and customized reporting detailing progress and achieved risk mitigations.

ModelOp

ModelOp focuses on life cycle management across traditional models, generative AI, agentic AI and third-party solutions. The platform integrates with AI development tools, data systems and infrastructure through over 50 out-of-the-box connectors. Governance capabilities include a centralized inventory that serves as an auditable system of record, supporting visibility into AI model usage, AI risk management, risk tiering and compliance-related activities. The platform facilitates policy enforcement through workflow automation and integrates with existing enterprise tools to support cross-functional processes. Reporting and monitoring features enable periodic reviews and risk assessments.

OneTrust

OneTrust supports privacy automation, consent and preference management, AI governance, data use governance, tech risk and compliance, and third-party risk management. The platform monitors vendor data, helps conduct risk assessments and supports regulatory compliance. Its Privacy Automation suite streamlines data mapping, subject requests and incident management through data discovery, AI-assisted workflows, interoperability and prebuilt templates. AI governance capabilities include automated registration and oversight of AI agents, datasets and models, with integration to AI development platforms and data repositories for life cycle management. The DataGuidance feature provides AI-powered search for regulatory content and side-by-side comparisons of privacy laws, supporting the interpretation of regulatory nuances.

Prodago

Prodago embeds regulatory, standards and policy requirements into AI projects, supporting automated governance planning and compliance. The SaaS platform uses Prodago's knowledge base to translate laws and internal policies into specific governance activities and artifacts, which can be adjusted based on project context and risk tolerance. The platform collects governance evidence and uses it to calculate AI trust level.

Relyance AI

Relyance AI combines AI governance and data security through a platform that continuously maps data flows from source code to AI models, enabling continuous visibility, automated compliance evidence and risk intelligence. The platform includes features like live model inventory, adaptive risk assessments and documentation to support regulatory compliance. Relyance AI also provides a visual Data Journeys interface for tracking data journeys, highlighting how and why data is processed across systems.

Saidot

Saidot's AI Governance Platform uses a knowledge graph to connect a company's AI inventory with governance items such as policies, risks and controls. The platform automates risk identification based on inherited risks from models and datasets, generates automated evaluation plans, suggests controls to manage identified issues, and supports ongoing monitoring. Saidot's governance platform also supports security evaluations and policy compliance with, for example, the EU AI Act. It also supports enforcement of documentation for use cases, risk classifications and acceptable use policies for AI assets like models and datasets

Securiti

Securiti's AI Security & Governance platform offers holistic oversight of AI systems across public clouds, private clouds and SaaS environments. Features include model discovery, risk classification, data mapping, data sanitization, runtime guardrails and data access controls. Compliance with over 20 global regulations and best practices frameworks is supported, where automated testing validates and provides evidence of prescribed controls, mitigating risks like sensitive data leakage, compliance violation and adversarial attacks. A centralized dashboard provides visibility and control, including the flow of sensitive data, enabling responsible and secure AI adoption across data ecosystems. See Note 4 for additional information.

ServiceNow

ServiceNow's AI Control Tower provides a centralized platform to govern, monitor and optimize AI systems. It supports bias detection, fairness audits and explainability features for any AI type. Governance workflows are customizable using playbooks, including legal, security and data reviews, integrated with ServiceNow's broader data model, workflow and case management tools. Strategy tools help align AI initiatives with enterprise goals, with roadmaps and portfolios. AI Control Tower supports the maintenance of a unified inventory of AI models, datasets and prompts throughout the AI life cycle. It provides dashboards to map AI use to business processes and track AI value, performance and compliance.

Truyo

Truyo supports the creation of an AI inventory by scanning websites, source code and content to detect AI use cases and assess associated risks. The platform provides structured assessments for both employee and vendor AI usage, generating a risk register highlighting policy creation and remediation workflows. The platform enables visibility into governance tasks, use cases and identified risks. It includes tools for testing and verifying bias and discrimination in AI models, and provides an audit log. Data remediation features allow companies to protect personal information while offering extended datasets for model testing and training. The product differentiates itself by including training for the company's workforce and AI governance board and by integrating AI consent and privacy automation.

Trustible AI

Trustible caters primarily to legal, risk, compliance and technology teams orchestrating the enterprise AI program. Its actionable insight layer delivers proprietary governance blueprints out of the box, including technical controls for developers and oversight teams. It centralizes AI use cases, agents, models, data and vendors in a unified inventory, for visibility, documentation and risk prioritization. Trustible includes customizable workflows and assessments supporting governance activities, such as reviewing AI proposals, assigning risk levels and tracking approvals. It supports model transparency with standardized documentation such as model cards, provides policy templates, risk and benefit taxonomies, and audit preparation tools for regulatory framework and internal standards alignment.

Market Recommendations

- Reassess your governance, risk management and compliance processes. Traditional processes may not effectively provide oversight of AI, and you may be required to partner with other governance, assurance and operations teams as well as the board of directors to adapt. It is critical to determine where shared control ownership is appropriate to avoid disputes and leadership turf wars that lessen the effectiveness of AI oversight.

- Decide what capabilities you require now. Some organizations prioritize investment in governance-oriented capabilities such as a catalog, risk management and regulatory policy engines. Other organizations will prefer to take a security-oriented approach by prioritizing runtime inspection and enforcement or AI usage control.

  - Capability mapping should be your primary foundation for engaging the right vendors by assessing how the AI governance platform meets the associated product criteria.

  - It is relevant to capture the different personas' expectations that relate to the agreed AI governance objectives to ensure the user experience and workflow capabilities can be provided by the selected technology and that the technology addresses a gap in your organization's technology stack.

- Prioritize whether your organization prefers to work with legacy vendors or if they are open to working with a startup.

  - Legacy vendors may have better financial soundness and a footprint in the organization's tech stack, but as vendors grow wide they often flatten product capabilities making them less useful for the end user.

  - Startups may better solve an organization's pain points but carry risks due to founder exit strategy, acquisitions and private equity/venture capital investments.

- Consider whether you prefer to invest in new technology or address risk monitored in disparate technologies using a business intelligence platform

Evidence

[1] NIST AI Risk Management Framework, U.S Government.

[2] ISO/IEC 42001:2023, ISO.

[3] EU AI Act: First Regulation on Artificial Intelligence, European Parliament.

[4] **2025 Gartner CEO and Senior Business Executive Survey**. This survey was conducted to examine CEO and senior business executive views on current business issues, as well as some areas of technology agenda impact. The survey was conducted from June 2024 through November 2024, with questions about the period from 2024 through 2026. One-quarter of the survey sample was collected from June through July 2024, and three-quarters was collected from October through November 2024. In total, 456 actively employed CEOs and other senior executive business leaders qualified and participated. The research was collected via 421 online surveys and 35 telephone interviews. The sample mix by role was CEOs (n = 303); CFOs (n = 95); COOs or other C-level executives (n = 39); and chairs, presidents or board directors (n = 19). The sample mix by location was North America (n = 194), Europe (n = 118), Asia/Pacific (n = 91), Latin America (n = 35), the Middle East (n = 15) and South Africa (n = 2). The sample mix by size was $50 million to less than $250 million (n = 32), $250 million to less than $1 billion (n = 122), $1 billion to less than $10 billion (n = 200) and $10 billion or more (n = 102). Disclaimer: Results of this survey do not represent global findings or the market as a whole, but reflect the sentiments of the respondents and companies surveyed.

[5] **2025 Gartner CIO and Technology Executive Survey.** This survey tracked how senior IT leaders worldwide prioritize strategic business, technical and management objectives. It was conducted online from 1 May through 28 June 2024. The survey includes respondents who lead an IT function, with a total of 3,186 CIOs and technology executives participating. The survey participants are representative of various geographies, revenue bands and industry sectors, including both public and private organizations. Disclaimer: The results of the survey do not represent global findings or the market as a whole, but reflect the sentiments of the respondents and companies surveyed.

[6] State of AI Agents, LangChain.

[7] **2024 Gartner Adapting to the New Risk Landscape Survey**. The survey was conducted to understand the role of data and analytics (D&A) leaders in mitigating risks related to AI and D&A, and how AI and D&A strategies and operating models can evolve to manage such risks more effectively. It explored cross-practice roles in AI and D&A risk management, the characteristics and effectiveness of risk practices for AI and D&A use cases, and the effects of risk management on an organization's ability to achieve business impact and technology adoption. The survey was conducted online from May through July 2024 among 387 respondents in North America (n = 231), EMEA (n = 90), Asia/Pacific (n = 59) and Latin America (n = 7). Qualified respondents were at the director level or above, with D&A as the primary focus of their work and, at most, one layer away from the highest-level D&A leader in their organization. Disclaimer: The results of this survey do not represent global findings or the market as a whole, but reflect the sentiments of the respondents and companies surveyed.

[8] AI Governance Market Advances with Stronger Regulatory Frameworks and Ethical AI Adoption, Precedence Research.

Note 1: Gartner's Initial Market Coverage

This Market Guide provides Gartner's initial coverage of the market and focuses on the market definition, rationale for the market and market dynamics.

Note 2: Vendor Selection

The list of vendors in this Market Guide were selected via analyst discretion to represent the range of products that support capabilities mentioned in the Market Description. Gartner used client interactions, vendor briefings and analyst research to determine that these offerings represent examples observed in this market. The vendors selected do not imply an exhaustive list. It provides a quick reference to AIGPs that are mentioned in Gartner client inquiries for a technology buyer primarily interested in satisfying the governance, risk management and compliance requirements for the organization where securing the organization is a secondary feature.

Note 3: FairNow AI

In October 2025, AuditBoard announced it has signed a definitive agreement to acquire FairNow AI and integrate it into AuditBoard's GRC platform. See AuditBoard Announces Acquisition of FairNow to Expand AI Governance Capabilities.

Note 4: Securiti

In October 2025, Veeam announced it has signed a definitive agreement to acquire Securiti AI. See Veeam Acquires Data Security Company Securiti AI for $1.7B.

Recommended by the Authors

Some documents may not be available as part of your current Gartner subscription.

Toolkit: Setting Up a Committee Charter for AI Governance

A Journey Guide to Manage AI Governance, Trust, Risk and Security

© 2026 Gartner, Inc. and/or its affiliates. All rights reserved. Gartner is a registered trademark of Gartner, Inc. and its affiliates. This publication may not be reproduced or distributed in any form without Gartner's prior written permission. It consists of the opinions of Gartner's Business and Technology Insights Organization, which should not be construed as statements of fact. While the information contained in this publication has been obtained from sources believed to be reliable, Gartner disclaims all warranties as to the accuracy, completeness or adequacy of such information. Although Gartner insights may address legal and financial issues, Gartner does not provide legal or investment advice and its insights should not be construed or used as such. Your access and use of this publication are governed by Gartner's Usage Policy. Gartner prides itself on its reputation for independence and objectivity. Its insights is produced independently by its Business and Technology Insights Organization without input or influence from any third party. For further information, see " Guiding Principles on Independence and Objectivity." Gartner insights may not be used as input into or for the training or development of generative artificial intelligence, machine learning, algorithms, software, or related technologies.

Table 1: Representative Vendors

Vendor	Product Name	Head quarters	Year Founded
Airia	Airia AI Security and Governance	Atlanta, Georgia, United States of America	2023
Cranium	Cranium AI Governance Platform	Short Hills, New Jersey, United States of America	2023
Credo AI	Credo AI Governance Platform	Palo Alto, California, United States of America	2020
Domino Data Lab	Domino Governance	San Francisco, California, United States of America	2014
Dynamo AI	DynamoGuard, DynamoEval	San Francisco, California, United States of America	2021
Fairnow	FairNow AI	McLean, Virginia, United States of America	2023
Holistic AI	Holistic AI	San Jose, California, United States of America	2020
IBM	watsonx.governance	Armonk, New York, United States of America	1911
LatticeFlow AI	AI GO!, AI Insights	Zurich, Switzerland	2020

Bahar Ghahremani bahar@newworldangels.com Mon Mar 16 10:12:42 2026

Vendor	Product Name	Head quarters	Year Founded
LogicGate	Risk Cloud	Chicago, Illinois, United States of America	2015
Monitaur	Monitaur Governance Platform	Boston, Massachusetts, United States of America	2019
ModelOp	ModelOp Center	Chicago, Illinois, United States of America	2018
OneTrust	OneTrust AI Governance	Atlanta, Georgia, United States of America	2016
Prodago	Prodago AI Governance	Montreal, Quebec, Canada	2013
Relyance AI	AI Governance Expert	San Francisco, California, United States of America	2020
Saidot	Saidot AI Governance Platform	Helsinki, Finland	2018
Securiti	AI Security & Governance	San Jose, California, United States of America	2019
ServiceNow	AI Control Tower	Santa Clara, California, United States of America	2004
Trustible	Trustible Responsible AI Governance Platform	Arlington, Virginia, United States of America	2023

Bahar Ghahremani bahar@newworldangels.com Mon Mar 16 10:12:42 2026

Vendor	Product Name	Head quarters	Year Founded
Truyo	AI Governance and Privacy Automation Platform	Chandler, Arizona, United States of America	2016

Source: Gartner

Bahar Ghahremani bahar@newworldangels.com Mon Mar 16 10:12:42 2026