

NWAI Investment Intelligence — AI Deal Pipeline

A visual reference for how Claude and the `nwai-tech-pipeline` plugin move a deal from raw pitch to IC-ready memo — and how the same platform scales across all six NWAI investment groups.

v0.22.0 · May 2026 · TechGroup



The Bigger Picture — Six Groups, One Platform

TechGroup is the first live instantiation. The architecture is designed to scale across all six NWAI investment groups.

Shared rails. Group-specific playbooks.


We are building an **investment intelligence platform** with common capabilities across the full dealflow, applied to the unique domain and investing thesis of each group. The pipeline stages, AI agents, Dealum CRM, and output formats are identical across all groups. What changes per group is the **investment playbook**: the screening gates that define what a qualified deal looks like, the scoring rubrics calibrated to that domain's risk profile, and the SME pool who bring the right domain expertise.

 TechGroup ✓ LIVE Full pipeline active. 5 themes. Dealum tag: "Tech".	 MedicalGroup PLANNED Needs: gates doc, SME table, thesis rubric, Dealum tag.	 SpaceGroup PLANNED Needs: gates doc, SME table, thesis rubric, Dealum tag.	 ConsumerGroup PLANNED Needs: gates doc, SME table, thesis rubric, Dealum tag.	 IndustrialGroup PLANNED Needs: gates doc, SME table, thesis rubric, Dealum tag.	 FintechGroup PLANNED Needs: gates doc, SME table, thesis rubric, Dealum tag.
--	--	--	---	---	--

- Shared across every group — built once - Dealum CRM integration (MCP server) - Claude orchestration layer + session context 8 AI agents: pipeline-monitor, company-researcher (PMTF), market-analyst, competitive-intelligence, technical-diligence, risk-assessor, plus the 3-agent financial diligence team — pricing-analyst → forecasting-analyst → venture-analyst - Pipeline stages: Inbox → Screen → Scout → Diligence → DD Report → Decision → Memo - Output formats: Triage Report, Scout Report, DD Kickoff Package, Post-Meeting Reconciliation, DD Investment Report, Exec Memo - Slash commands: <code>/sync-pipeline</code> <code>/screen</code> <code>/scout</code> <code>/diligence</code> <code>/post-meeting</code> <code>/dd-report</code> <code>/decision</code> <code>/memo</code>	- Group-specific — the playbook per group - Dealum tag filter (e.g., "Tech", "Medical", "Space") - Screening gates & flags — what constitutes a qualified deal in this domain - Scoring rubrics — moat & risk dimensions calibrated to domain risk profile - Investment themes — domain sub-categories for deal mapping - SME pool & lead assignment table — the domain experts who validate deals - Investment thesis — the codified investability logic unique to each group
--	---


1 The Pipeline — Seven Stages

Every Tech-tagged deal in Dealum moves left to right. Claude tracks the current stage and calibrates analysis depth accordingly.



2 Commands — What to Type

Type these in the Cowork chat window. Claude follows the corresponding workflow and updates Dealum automatically.



AI Agents — Claude's Research Team

- 3 Eight specialized agents. The Stage 2A parallel set runs first (~5–8 min); the Stage 2B financial diligence chain runs sequentially because each downstream agent consumes upstream output. Each agent reports back with a structured briefing that feeds the relevant command output.

pipeline-monitor

Triggered by: session start, /sync-pipeline, "what's in the pipeline?"

Connects to Dealum, pulls all Tech-tagged deals, groups them by pipeline stage, and surfaces stalled deals with suggested next actions.

Returns: Live pipeline dashboard — deal counts per stage, attention items, recommended next command.

company-researcher **PMTF refresh · v2.13**

Triggered by: /scout (full) + /diligence (Stage 2A parallel)

Team-first mandate. **Founder Claim Verification Protocol** (LinkedIn + Perplexity cross-reference of every specific exit/ARR/role claim →  VERIFIED /  PARTIAL /  UNVERIFIED /  CONTRADICTED). **Team-Level PMTF** (skills coverage across domain depth / engineering / market access; explicit gap detection). **Team Commitment Depth** (full-time vs advisor ratio, "stuck" advisor flagging). Plus supporting context: funding, traction, commercial validation, public red flags.

Returns: People Verification Brief + Team Quality + PMTF Assessment + Commitment Depth + commercial briefing → feeds Team Validation + Commercial Validation (Layer 2).

market-analyst

Triggered by: /scout (full) + /diligence (full)

Structural discontinuity test, independent TAM/SAM/SOM validation, market timing scoring, industry dynamics, and customer behavior signals.

Returns: Market Analysis Briefing → feeds Market Validation (Layer 2).

competitive-intelligence

Triggered by: /scout (full) + /diligence (full)

Maps direct competitors and incumbent threats, alternative solutions, strategic incumbents, positioning analysis, and moat input data.

Returns: Competitive Intelligence Briefing → feeds Competitive Validation (Layer 2).

technical-diligence

Triggered by: /scout (light — thin wrapper + TRL) + /diligence (full)

Thin wrapper detection (3 tests), TRL assessment, IP/patent research, technical architecture signals, AI moat input data, technical team depth.

Returns: Technical Diligence Briefing → feeds Technical Validation (Layer 2).

risk-assessor

Triggered by: /scout (light — top kill risks) + /diligence (Stage 2A parallel)

Regulatory risk research, exit landscape + acquirer dynamics, execution risk signals, market and financial risk signals from public sources.

Returns: Risk Assessment Briefing → feeds cross-cutting risk overlay across all Layer 2 groups.

STAGE 2B — 3-AGENT FINANCIAL DILIGENCE TEAM **NEW V2.13**

Sequential chain: **pricing-analyst** → **forecasting-analyst** → **venture-analyst**. Replaces the legacy financial-analyst agent. Built on the McMurry method and the "no AI slop" rule: build the proprietary forecast first, then compare to the company's submitted numbers — never just render the founder's deck.

\$ pricing-analyst **NEW · v2.13**

Triggered by: /diligence (Stage 2A parallel)

Pricing Maturity score (PROVEN / EARLY SIGNALS / DISCOVERY / UNKNOWN). Value-proposition validation with quantified pricing-to-value ratio. **Channel Economics** 5-year margin compression model (third-party distributors typically command a meaningful markup multiple of the direct price; this gap compresses over time as the channel partner gains volume leverage). Unit economics (CAC / LTV / payback). Pricing Pressure Forecast (commoditization risk).

Returns: Pricing Analyst Briefing with explicit "Feed Forward to Forecasting Analyst" section → feeds Financial Validation (Layer 2) + DD Report S5 (GTM).

📊 forecasting-analyst **NEW · v2.13**

Triggered by: /diligence (Stage 2B sequential — after pricing-analyst)

Independent 5-year financial forecast using the **McMurry method**: build the proprietary forecast from comps + pricing analysis, THEN compare to the company's submitted numbers. Enforces the "no AI slop" rule — every Bear/Base/Bull line carries a mandatory *because* clause. Outputs P&L, cash flow, balance sheet, capital plan with round timing/sizing, and a **Founder Financial Literacy Assessment**.

Returns: Independent forecast with comparison delta to company's submitted model → feeds Financial Validation (Layer 2) + DD Report S9 (Financials).

🏠 venture-analyst **NEW · v2.13**

Triggered by: /diligence (Stage 2B sequential — after forecasting-analyst)

Final synthesis layer. Defensible valuation today (3 methods reconciled). Projected exit Y3 + Y5 (Bear/Base/Bull). Capital + dilution model. **IRR / multiple / 35% IRR hurdle test**. **NWA 10x-in-5-years criterion** check (pass/fail with math). Deal structure recommendation: priced equity / convertible with cap / participating preferred / re-priced terms / decline. Closes the founder-claimed-vs-defensible-analysis gap with documented methodology.

Returns: Valuation conclusion + structural recommendation → feeds DD Report S8 (Deal Structure), S9 (Financials), S11 (Exit) + /decision + /memo Slide 4.

4

Investment Framework — The Playbook

Reference documents Claude loads when running analysis. These encode NWAi's proprietary criteria so the same standards are applied consistently, regardless of how compelling a pitch sounds.

Reference File	What It Contains	Used At
<code>gates-and-flags.md</code>	NWAi Universal Triage Framework v2.0 (all verticals): 3 hard gates + NWA Filter (Cynical Default, Goliath Test, LLM Ingestion Test, Revenue Quality Audit) + Opportunity 6x0–5 (/30) + Readiness 5x0–5 (/25) + 3-search Research Protocol + decision logic (ADVANCE ≥ 20/30, WATCH 14–19, DECLINE < 14)	Screen (universal)
<code>gates-and-flags-techgroup.md</code>	TechGroup extension: Track A (Software/AI/Cloud) and Track B (Hardware/Robotics) 6-dimension Opportunity rubrics, AI Wrapper Assessment, Replicability Speed Flag, Hardware Last Mile Standard, TRL Hard Cap	Screen (TechGroup)
<code>_archive/gates-and-flags-6gate-legacy.md</code>	Pre-v2.0 legacy 6-gate pass/fail AutoKill framework — superseded by Universal Triage in April 2026; retained for reference only	Reference only
<code>scout-questions.md</code>	Phase 1 viability (Big Idea / Market / Moat) + Phase 2 execution (Team with PMTF + Commitment Depth + verification sub-assessments, Technology, Traction, GTM, Exit) + Scout Conviction Score	Scout
<code>diligence-scoring-rubrics.md</code>	General Moat (0–6), AI Moat (0–10), Risk (1–10) with explicit numeric criteria for team commitment ratio + PMTF score + founder claim verification, Bear/Base/Bull model with mandatory <i>because</i> clauses + time-varying ARPU (channel pressure modeling)	Diligence
<code>dd-checklist.md</code>	17-folder due diligence framework with team assignment slots and gate-critical folder designations	Diligence
<code>ai-moats-framework.md</code>	Three moat types for AI defensibility + Replicability Speed Matrix (4-row threat actor table with Yellow Flag triggers)	Scout + DD Report + Memo
<code>dd-report-format-reference.md</code>	Exact layout spec for the DD Report: 11 sections, 2-column table format, RAG score colors, content rules	DD Report
<code>memo-format-reference.md</code>	4-slide PPTX structure, NWAi branding spec, ++/-- notation, "What is the Bet?" framing	Memo
<code>diligence-analysis-framework.md</code>	Analyst lens behavioral standard for post-meeting analysis: dual role (analyst POV primary / tracker secondary), declaration taxonomy, meeting-type playbooks, progressive POV architecture	Diligence (post-meeting)

TWO ANALYTICAL LENSES — ALWAYS ACTIVE, AT EVERY STAGE

⚡ **Structural Discontinuity**

Is this Big Idea riding a genuine, irreversible market shift — or incremental improvement dressed up as disruption? A real discontinuity changes the rules of the game. An incremental product just plays the existing game a little better.

🔒 **Memory Lock-in**

Does the product get smarter, stickier, or more embedded the more it's used? Proprietary data flywheels, deep workflow integration, and switching costs driven by accumulated context are the hallmarks of a durable AI moat.

DD REPORT SCORING — 11 SECTIONS, 1–5 SCALE, RAG COLORED

Score 4–5 → Green

■ Strong signal, advance confidently

Score 3 → Amber

■ Conditional — specific concerns to resolve

Score 1–2 → Red

■ Serious concern — likely Kill condition

5 Non-Negotiable Investment Criteria

Any deal that fails any of these is a Pass. No exceptions — regardless of how compelling the pitch or team appears.



Structure

Equity required. Priced round preferred (Seed–Series A). No SAFEs. Convertible debt acceptable. C-Corp only.



Geography

US-based HQ and IP. Executive team must be US-based. No foreign-owned entities.



Traction

Real revenue or signed paying customers. Research projects and demo labs do not qualify.



Stage

MVP complete with at least one successful beta completed. TRL 5+ (GAO scale).



Syndication

Credible lead investor required. Cap table must be clean — no structural red flags.



Return Potential

10x in 5 years minimum. TAM must support a 20x–100x return path.

6 What Gets Saved — The Output Chain

Each command saves its output to the `deals/` folder. Downstream commands load the most recent prior-stage file automatically — so context carries forward without manual copy-paste.

`/screen`



Triage Report (Word doc) — gates verdict, opportunity/readiness scores, live pitch questions
[Company] - Triage Report YYYY-MM-DD.docx

`/scout`



Scout Assessment Report (Word doc) — one-sentence verdict, biggest risk, 3–5 diligence questions, theme mapping
[Company] - Scout Assessment Report YYYY-MM-DD.docx

`/diligence`



DD Kickoff Package (Word doc) — scored rubrics, 17-folder assignments, agent research briefings
[Company] - DD Kickoff Package YYYY-MM-DD.docx

`/post-meeting`



Post-Meeting Reconciliation (Word doc) — Analyst POV + Key Insights table + Resolved/Open tracker
[Company]-[Type]-Meeting-Reconciliation-YYYY-MM-DD.docx

`/dd-report`



DD Investment Report (Word doc) — 11 scored sections, RAG table layout, IC-ready
[Company]-NWAi-DD-Report-YYYY-MM-DD.docx

`/decision`



Dealum updated — pipeline step advanced, decision tag applied (Invest / Pass / Watch)
No file output — CRM action only

`/memo`



Executive Summary (4-slide PPTX) — NWAi branding, IC presentation format
[Company]-NWAi-Exec-Summary-YYYY-MM-DD.pptx