

TALKING-POINTS CRIB SHEET

"What I want to walk you through over lunch is four things — what we've already built, what it takes to evolve it, how we keep it enterprise-grade as members come on, and the bigger prize underneath it all."

01

What we've built — the TechGroup deal pipeline

VERDICT

Built the analytical layer Dealum doesn't have. Dealum stays authoritative for intake, stage tracking, member records, and decisions logged — the pipeline I built generates the intelligence at each stage (screening verdicts, scout assessments, diligence packages, DD Reports, Memos). End-to-end, processing real deals on my desktop today.

- Six-layer system: filesystem deal data → 7-stage pipeline (Inbox → Memo) → 8 slash commands → 9 AI research agents → 10 proprietary framework documents → persistent session context.
- Agents do real work: company research, market analysis, competitive intel, technical diligence, risk, pricing, independent 5-yr financial forecast (McMurry method), valuation with 35% IRR hurdle test.
- Outputs match our visual standard — DD Reports (11 scored sections, RAG-coded) and 4-slide Memos rendered exactly to the STL/Synergist reference format.
- Filesystem-first today; Dealum API integration is deferred. Runs across Claude Desktop, Claude Code CLI, and Cowork.
- **Limits to name explicitly:** Single-user (me). One deal at a time. TechGroup only. Not on the web. Not multi-tenant. Not member-accessible.

SAY IT THIS WAY

"I've proven the rails work. The same pattern that screens, scouts, and runs diligence for TechGroup is what we lift onto the platform — but the platform is where it stops being a single-person tool."

What we need to build — multi-user, multi-group, web

VERDICT

The build is two things at once: lift the proven desktop pipeline onto a multi-tenant web platform AND build the member intelligence substrate the pipeline plugs into. V1 ships the minimum substrate in 30–45 days; full 6-group rollout in 3–4 months; substrate before activations.

Layer	What it owns during V1–V2
Dealum (stays authoritative)	Deal intake · stage tracking · member records · decisions logged · historical archive
Platform — Deal Intelligence (V1)	The 9-agent pipeline lifted onto the web: screening verdicts, scout assessments, diligence packages, DD Reports, Memos
Platform — Member Intelligence (V1)	The substrate: member graph · decisions ledger · Match Engine · SME POV capture · social deal card

Beyond V3 the platform subsumes Dealum's roles too — but that's the long arc, not the V1 build.

- **V1 (30–45 days): two tracks at once** — deal intel (the 9 agents on the web) + member intel substrate (graph, ledger, Match Engine, social deal card).
- **V2 (~3 months total): the other 5 groups.** Medical → Space → Consumer → Industrial → Fintech, 2 weeks each. Each plugs its own playbook (rubrics, themes, SMEs) into universal rails. Adding a group is config, not code.
- **V3+: ecosystem sourcing.** Proactive feeds from incubators, accelerators, universities. Pattern matching upstream of Track 1.
- **Beyond V3:** native intake + founder portal + native CRM → Dealum decommissioned. Full platform parity.
- Architected from V1 for 6 groups even though only TechGroup is live — the data model survives every later transition.

SAY IT THIS WAY

"V1 is 30–45 days for TechGroup. All six groups within 3–4 months. The total arc to full parity — including replacing Dealum — is 12–18 months. We don't compress that; we sequence it."

Enterprise-grade foundation — the desktop-to-platform transition

VERDICT

Enterprise-grade isn't a security checkbox — it's the transition from a session-dependent desktop tool that processes one deal at a time to an always-on platform that serves 200+ members across 6 groups. Four dimensions define it: always-on reliability, multi-user access, clean Dealum integration, and secure data handling.

Capability	Desktop today (Cowork / Claude Code)	Enterprise platform (target)
Users	Me — single session	200+ members with role-based access
Reliability	Session-dependent — runs when I'm working	Always-on cloud — processes deals as they enter Dealum
Intelligence delivery	I manually share outputs	Automated push: right output to right member at each stage
SME input	Ad-hoc email + conversation	Structured form integrated with the diligence pipeline
Dealum integration	I read Dealum, update manually	Bidirectional — automated ingestion + write-back
Security	Desktop session, no formal controls	Encrypted, RBAC (6 × 6), append-only audit, member data rights
Scale	Single thread, one deal at a time	Full deal flow (1,000+/year) concurrently

- **Always-on reliability:** Intelligence runs on a cloud server, not my Cowork or Claude Code session. A deal lands in Dealum → platform processes it automatically. Not session-dependent.
- **Multi-user access:** Reports and deal cards delivered to (or accessible by) the right members at the right stage. SMEs submit diligence input via a structured surface, not email. Role-based: 6 groups × 6 roles · per-group AND per-deal scoping · cross-group opt-in.
- **Clean Dealum integration:** Bidirectional. Dealum stays authoritative for deal tracking; the platform owns analysis. Read in, write back. (Long-term — Beyond V3 per Section 2 — the platform subsumes Dealum.)
- **Secure data handling:** Encrypted in transit + at rest. Append-only decisions ledger as the audit trail. GDPR-equivalent member data rights from day one. Vertical overlays (HIPAA-adjacent for Medical, regulatory for Fintech/Space) layer on top.

SAY IT THIS WAY

"What we're building isn't a desktop AI tool that's been scaled up — it's a real platform. The shift is from me processing one deal at a time on my Mac to an always-on system delivering intelligence to the right member at the right stage. Security is the floor; reliability, multi-user access, and bidirectional Dealum integration are what actually makes it enterprise."

From deal intelligence → member/social intelligence platform

VERDICT

The deal pipeline is the wedge. The enduring asset is the substrate — member expertise graph plus 20+ years of indexed IC reasoning. AI analytical capability is commoditizing fast; what no vendor can sell us and no competitor can copy is the network and the decision memory compounded over time.

- **The substrate has five components:** Sensors (what writes), Member Graph (who knows what), Decisions Ledger (append-only IC history), Policy Layer (NWAi criteria as code), Learning Loop (every cycle re-tunes the next).
- **The thesis (from my May 23 to you):** AI compresses screening, scaffolding, DD synthesis. It doesn't reach selection under uncertainty, conviction-forming, or the relationship side of value-add. **The work AI compresses isn't the work that produces returns.**
- **Now / near / far:**
 - *Now:* Build deal intelligence as the interim operational capability AND start the member/social layer in parallel. Two builds, integrated by design.
 - *Near (12–24 mo):* Analytical capability commoditizes. Swap modules at market price; reinvest into substrate depth.
 - *Far:* Substrate is the moat; AI is the connective tissue activating it. Analytical layer = commodity input. Network + institutional memory = what compounds.
- **What members actually want:** Network, founders, energy, SME engagement. Returns matter, but they're the 4th of 5 motivators. The platform has to serve that — not just process deals faster.

SAY IT THIS WAY

"The deal pipeline is real and it works. But if all we build is faster diligence, we've built a tool a vendor will replicate in 18 months. The compounding asset is the member graph and the decisions ledger. That's what we're really building."

CLOSING — THE DEALUM FRAMING + THE ASK

ON DEALUM

"Near-term, Dealum keeps the firehose handled — which lets us focus the build on what's actually missing: deal intelligence and member intelligence. Over time, the platform encompasses everything Dealum does today — native intake, founder portal, member roster, CRM — and Dealum is decommissioned. That's a Beyond-V3 conversation, not a V1 one. But every architectural decision we make in V1 anticipates it."

THE SINGLE BIGGEST RISK (NAME IT BEFORE THEY DO)

"AI erodes the high-ranking experiential motivators — peer networking, founder proximity, SME engagement — broadly enough that members disengage before the substrate compounds. That's why substrate ships in V1, not later."

THE ASK (FOR THE INCOMING DEVELOPER-MEMBER)

"You join NWA as an advisory member while building this. You're inside the substrate you're building. V0 is 2–4 weeks of architecture proposal; V1 SOW follows. The arc through V2 is ~4 months from kickoff; full parity is 12–18 months."

Source docs: GUT v1.0 [docs/strategy/foundational/](#) · Enterprise Build Brief v1.1 [docs/reference/NWA Investment Intelligence Platform/](#) · Member Intelligence V1 Features [docs/strategy/foundational/](#) · Architecture v0.26.0 (workspace root)
· May 23 email to Ron [docs/strategy/member-communication/](#) .